

QUEEN MARY, UNIVERSITY OF LONDON
SCHOOL OF ELECTRONIC ENGINEERING AND COMPUTER SCIENCE

ECS708U: Machine Learning

Assignment 1 (Part 1) - Linear Regression

Mughees Asif | 180288337

Due: 12 October 2021

1 Linear Regression with One Variable

Calculate the hypothesis for the i^{th} sample of X , given X , θ and i :

```
1 # The hypothesis for the i-th sample of X, given X, theta and i, is the
  ↳ dot product of for each sample with theta can be calculated as:
2
3 hypothesis = np.dot(X[i], theta)
4
```



Modified calculate_hypothesis function:

```
1 # The `calculate_hypothesis` method can be called with X, theta, and i:
2 hypothesis = calculate_hypothesis(X, theta, i)
3
```

Effect of the learning rate α :



- $\alpha = 1e^{-6}$

Minimum cost: 31.85851. on iteration #50





- $\alpha = 1e^{-4}$

Minimum cost: 17.36882, on iteration #50





- $\alpha = 1e^{-2}$

Minimum cost: 5.67829, on iteration #50





- $\alpha = 0$

Minimum cost: 32.07273, on iteration #1



- $\alpha = 1$

Minimum cost: 172570.09522, on iteration #1



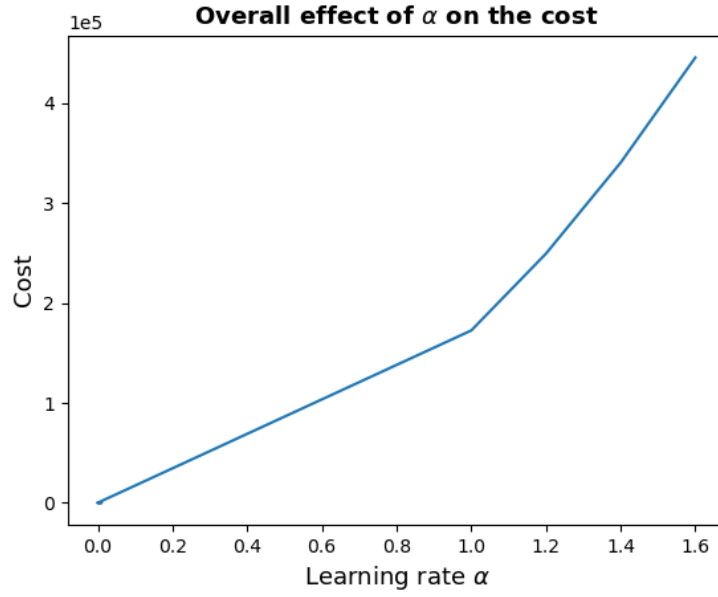


Figure 6: Learning rate and the cost

The learning rate hyperparameter α determines the amount to update the weights during gradient descent to minimise the back-propagation error loss. When α is too small ($1e^{-6}$), the **prediction** does not correlate with the **groundtruth** as the model did not have enough iterations to complete the minimisation or only reached the local minima. However, when $\alpha = 1.6$, the model overshoots (effectively missing the global maxima) and thereby, a vast increase in the cost is also incurred.

From Fig. 6, it can be observed that as α increases from a very small value to a larger value, the gradient is stable when $0 < \alpha < 1$, where a sudden increase in the cost is observable when $\alpha > 1$. The **prediction** is closer to the **groundtruth** when $\alpha = 1e^{-2}$ with a minimal cost value. Therefore, an optimum α for this model exists between the interval $[0, 1]$.

2 Linear Regression with Multiple Variables

Supporting the new hypothesis function:



```
1 # The hypothesis for the i-th sample of X, given X, theta and i, is the  
  ↪ dot product of for each sample with theta can be calculated as:  
2 hypothesis = np.dot(X[i], theta)  
3
```

```
1 # The `calculate_hypothesis` method can be called with X, theta, and  
  ↪ i:  
2 hypothesis = calculate_hypothesis(X, theta, i)  
3  
4 output = y[i]  
5  
6 # Update `sigma`: keep the summation for the derivative term  
7 sigma = sigma + (hypothesis - output) * X[i]  
8  
9 # Update `theta_temp`  
10 theta_temp = theta_temp - (alpha / m) * sigma  
11
```

Effect of the learning rate α :



- $\alpha = 1e^{-2}$

Minimum cost: 10596969344.16698, on iteration #100
Theta: [215810.61679138 61446.18781361 20070.13313796]





- $\alpha = 0$

Minimum cost: 65591548106.45744, on iteration #1
Theta: [0. 0. 0.]





- $\alpha = 1$

Minimum cost: 2043280050.60283, on iteration #48
Theta: [340412.65957447 109447.79646964 -6578.35485416]





- $\alpha = 1.2$

Minimum cost: 2043280050.60802, on iteration #100
 Theta: [340412.65957447 109447.73880033 -6578.41252348]





- $\alpha = 1.4$

Minimum cost: 11667746227.66517, on iteration #2
 Theta: [3.40412659e+05 -1.10851714e+12 -1.10851726e+12]



Evidently, the optimum value for α exists when $0 < \alpha < 1$, therefore, after experimenting with several values in the interval $[0, 1]$, 0.1 was identified as a suitable trade-off between convergence and overshooting:

Minimum cost: 2043462824.61817, on iteration #100
 Theta: [340403.61773803 108803.37852266 -5933.9413402]



Making a prediction:

```
1 # Create two new samples: (1650, 3) and (3000, 4)
2 houses = np.array([[1650, 3],
3                   [3000, 4]])
4 # Make sure to apply the same preprocessing that was applied to the
  ↪ training data
5 houses, mean_vec_new, std_vec_new = normalize_features(houses)
6 # After normalizing, we append a column of ones to X, as the bias term
7 column_of_ones = np.ones((houses.shape[0], 1))
8 houses = np.append(column_of_ones, houses, axis=1)
9 # Calculate the hypothesis for each sample, using the trained parameters
  ↪ theta_final
10 prices = []
11 for i in range(len(houses)):
12     prices.append(calculate_hypothesis(houses, theta_final, i))
13
14 # Print the predicted prices for the two samples
15 print("A house with 1650 sq. ft. and 3 bedrooms will cost:
  ↪ {:.5f}".format(prices[0]))
16 print("A house with 3000 sq. ft. and 4 bedrooms will cost:
  ↪ {:.5f}".format(prices[1]))
17
```

Minimum cost: 2043462824.61817, on iteration #100

Theta: [340403.61773803 108803.37852266 -5933.9413402]

A house with 1650 sq. ft. and 3 bedrooms will cost: 237534.18056

A house with 3000 sq. ft. and 4 bedrooms will cost: 443273.05492

3 Regularized Linear Regression



gradient_descent method modifications:

```
1      # The `calculate_hypothesis` method can be called with X, theta, and  
2      ↪ i:  
3      hypothesis = calculate_hypothesis(X, theta, i)  
4      output = y[i]  
5      sigma = sigma + (hypothesis - output) * X[i]  
6  
7      # update theta_temp  
8      bias = theta_temp[0]  
9      theta_temp = theta_temp * (1 - (alpha * (1 / m))) - (alpha / m) * sigma  
10     theta_temp[0] = bias - (alpha / m) * sigma[0]  
11  
12     # copy theta_temp to theta  
13     theta = theta_temp.copy()  
14  
15     # append current iteration's cost to cost_vector  
16     iteration_cost = compute_cost_regularised(X, y, theta, 1)
```

Effect of the learning rate α , and the regularization parameter λ :



- The optimum value for α was found to be **0.02**.
- $\alpha = 0.02$, $\lambda = 0.0$

Minimum cost: 0.01215, on iteration #200





- $\alpha = 0.02$, $\lambda = 0.5$

Minimum cost: 0.03472, on iteration #200




- $\alpha = 0.02$, $\lambda = 1$

Minimum cost: 0.05370, on iteration #200



- $\alpha = 0.02$, $\lambda = 5$

Minimum cost: 0.13857, on iteration #200



The regularization parameter λ is a scalar value that reduces overfitting. From the above plots, when $0 < \lambda < 1$, similar to α , the **prediction** is closest to the **groundtruth**. As λ is increased, the rate of the gradient descent also increases. Consequently, the minimum cost reduces and the model fits the data suitably. However, increasing the parameter beyond a threshold (≈ 1) results in underfitting.