

QUEEN MARY, UNIVERSITY OF LONDON
SCHOOL OF ELECTRONIC ENGINEERING AND COMPUTER SCIENCE

ECS708U: Machine Learning


Assignment 2: Clustering and Mixture of Gaussians (MoG)

Mughees Asif | 180288337


Due: 12 October 2021

Task 1

```
1 # Store f1 in the first column of X_full, and f2 in the second column of
  ↪ X_full
2 X_full[:, 0] = f1
3 X_full[:, 1] = f2
4 ...
5 # Create array containing only samples that belong to phoneme 1
6 X_phoneme_1 = X_full[phoneme_id == p_id, :]
7
```



(a) Phoneme 1



(b) All phonemes

Figure 1: Fundamental frequencies, f_1 against f_2

f1 statistics:

Min: 190.00 Mean: 563.30 Max: 1300.00 Std: 201.1881 | Shape: 1520

f2 statistics:

Min: 560.00 Mean: 1624.38 Max: 3610.00 Std: 636.8032 | Shape: 1520

Task 2

```
1      # Store f1 in the first column of X_full, and f2 in the second column of  
    ↪ X_full  
2      X_full[:, 0] = f1  
3      X_full[:, 1] = f2  
4      ...  
5      # Create array containing only samples that belong to phoneme 1  
6      X_phoneme_1 = X_full[phoneme_id == p_id, :]  
7
```

$k = 3$

First run




Figure 2: Gaussian phoneme, $k = 3$

Implemented GMM | Mean values

```
[ 270.3952 2285.4653]
```

```
[ 350.8446 3226.3394]
```

```
[ 312.59125 2783.898  ]
```

Implemented GMM | Covariances

```
[[ 1213.73843494    0.          ]
```

```
 [    0.          14278.42029955]]
```

```
[[ 4102.87537495    0.          ]
```

```
 [    0.          27829.54221723]]
```

```
[[3562.59743765    0.          ]
```

```
 [    0.          7657.84897216]]
```

Implemented GMM | Weights

```
[0.43514434 0.18386038 0.38099528]
```

Second run




Figure 3: Gaussian phoneme, $k = 3$

Implemented GMM | Mean values

```
[ 312.7931 2811.8486]
[ 362.2751 3270.6072]
[ 269.92792 2281.0254 ]
```

Implemented GMM | Covariances

```
[[ 3133.01971418    0.          ]
 [    0.          20592.59093668]]
[[ 5064.81166948    0.          ]
 [    0.          29330.22701612]]
[[ 1215.82904825    0.          ]
 [    0.          13740.86780726]]
```

Implemented GMM | Weights

```
[0.44128011 0.13448926 0.42423063]
```

$k = 6$

First run




Figure 4: Gaussian phoneme, $k = 6$

Implemented GMM | Mean values

```
[ 457.8011 3416.4023]
[ 247.18047 2210.3774 ]
[ 311.85373 2778.3267 ]
[ 301.0684 2344.1323]
[ 218.6318 2363.589 ]
[ 332.71387 3171.762 ]
```

Implemented GMM | Covariances

```
[[ 7474.90630756 0. ]
 [ 0. 16556.63598171]]
[[ 279.05520316 0. ]
 [ 0. 12839.83287966]]
[[3684.31355542 0. ]
 [ 0. 7070.6920406 ]]
[[ 421.65474188 0. ]
 [ 0. 7209.39225671]]
[[ 25.51633687 0. ]
 [ 0. 994.57566331]]
[[ 1251.69581206 0. ]
 [ 0. 27146.11552164]]
```

Implemented GMM | Weights

```
[0.02607926 0.19808257 0.36774705 0.20403745 0.03192522 0.17212846]
```

Second run




Figure 5: Gaussian phoneme, $k = 6$

Implemented GMM | Mean values

```
[ 259.19806 2777.2854 ]
[ 347.8378 2803.3938 ]
[ 285.69077 2701.1602 ]
[ 332.73532 3219.3318 ]
[ 451.44778 3101.4407 ]
[ 269.71927 2277.7375 ]
```

Implemented GMM | Covariances

```
[[1360.26534051 0. ]
 [ 0. 4540.09031381]]
[[ 845.34786211 0. ]
 [ 0. 9868.5388874 ]]
[[2.60575697e+01 0.00000000e+00]
 [0.00000000e+00 3.17874345e+04]]
[[ 1264.84075501 0. ]
 [ 0. 20058.38684489]]
[[ 5387.36721704 0. ]
 [ 0. 118675.67982747]]
[[ 1252.55905729 0. ]
 [ 0. 13264.79104545]]
```

Implemented GMM | Weights

```
[0.13982013 0.18456919 0.06574364 0.14693026 0.04560488 0.41733191]
```

Task 3

```
1 # Store f1 in the first column of X_full, and f2 in the second column of
  ↪ X_full
2 X_full[:, 0] = f1
3 X_full[:, 1] = f2
4 ...
5 # Create an array containing only samples that belong to phoneme 1 and
  ↪ samples that belong to phoneme 2
6 X_phonemes_1_2 = X_full[np.logical_or(phoneme_id == 1, phoneme_id == 2),
  ↪ :]
7 ...
8 # Get predictions on samples from a GMM with k components, pretrained on
  ↪ phoneme 1
9 phoneme_model_1 = 'data/GMM_params_phoneme_{:02}_k_{:02}.npy'.format(1,
  ↪ k)
10 params_1 = np.load(phoneme_model_1, allow_pickle=True).item()
11 copy = X_phonemes_1_2.copy()
12 Z_1 = get_predictions(
13     params_1['mu'],
14     params_1['s'],
15     params_1['p'],
16     copy
17 )
18 pred_1 = Z_1.sum(axis=1)
19
20 # Get predictions on phoneme 2
21 phoneme_model_2 = 'data/GMM_params_phoneme_{:02}_k_{:02}.npy'.format(2,
  ↪ k)
22 params_2 = np.load(phoneme_model_2, allow_pickle=True).item()
23 Z_2 = get_predictions(
24     params_2['mu'],
25     params_2['s'],
26     params_2['p'],
27     copy
28 )
29 pred_2 = Z_2.sum(axis=1)
30
31 # Compare these predictions for each sample of the dataset, and calculate
  ↪ the accuracy, and store it in a scalar variable named "accuracy"
32 predictions = np.ones(len(copy)) * 2
33 predictions[pred_1 >= pred_2] = 1
34 labels = phoneme_id[np.logical_or(phoneme_id == 1, phoneme_id == 2)]
35 accuracy = np.sum(predictions == labels) / copy.shape[0] * 100
36
```


Results

Accuracy using GMMs with 3 components: 96.38%
Accuracy using GMMs with 6 components: 95.72%

Observations

The above code is structured as:

1. **Line 2–4:** Copy the relevant frequency values.
2. **Line 6:** Fill `X_phonemes_1_2` with the samples of `X_full` that belong to the chosen phonemes.
3. **Line 9–10:** Load the pretrained GMM model for $k = 3$ & 6.
4. **Line 11:** Make a copy of the `X_phonemes_1_2` to get the samples for `p_id`.
5. **Line 12–17:** Get the predictions (likelihood) for both phonemes. Sum up the predictions and save in `pred_i`, $i = 1 \mid 2$.
6. **Line 21–29:** Repeat steps 3–5 for the second phoneme.
7. **Line 33–36:** Initialise the `predictions` array that holds the cumulative likelihood from the individually summed up predictions from step 5. Use the information to calculate the accuracy %.

The GMMs trained on 3 parameters were found to train faster and exhibit increased accuracy. The GMMs with 6 parameters were found to exhibit a reduction in accuracy due to overfitting and took longer to train.

Task 4

```
1      # Store f1 in the first column of X_full, and f2 in the second column of
      ↪ X_full
2      X_full[:, 0] = f1
3      X_full[:, 1] = f2
4      ...
5      # Create an array containing only samples that belong to phoneme 1 and
      ↪ samples that belong to phoneme 2
6      X_phonemes_1_2 = X_full[np.logical_or(phoneme_id == 1, phoneme_id == 2),
      ↪ :]
7      ...
8      # Generate grid
9      ax_f1 = np.linspace(min_f1, max_f1, N_f1)
10     ax_f2 = np.linspace(min_f2, max_f2, N_f2)
11     x_axis, y_axis = np.meshgrid(ax_f1, ax_f2)
12     samples = np.stack((x_axis.flatten(), y_axis.flatten())).transpose()
13
14
15     # Phoneme model no. 1
16     phoneme_model_1 = 'data/GMM_params_phoneme_{:02}_k_{:02}.npy'.format(1,
      ↪ k)
17     params_1 = np.load(phoneme_model_1, allow_pickle=True)
18     params_1 = np.ndarray.tolist(params_1)
19     copy = samples.copy()
20     Z_1 = get_predictions(
21         params_1['mu'],
22         params_1['s'],
23         params_1['p'],
24         copy
25     )
26
27     # Phoneme model no. 2
28     phoneme_model_2 = 'data/GMM_params_phoneme_{:02}_k_{:02}.npy'.format(2,
      ↪ k)
29     params_2 = np.load(phoneme_model_2, allow_pickle=True)
30     params_2 = np.ndarray.tolist(params_2)
31     Z_2 = get_predictions(
32         params_2['mu'],
33         params_2['s'],
34         params_2['p'],
35         copy
36     )
37
38     # Get the predictions
39     pred_1 = Z_1.sum(axis=1)
40     pred_2 = Z_2.sum(axis=1)
41     predictions = np.ones(len(copy)) * 2
42     predictions[pred_1 >= pred_2] = 1
43
```

```

44 # M
45 M = predictions.reshape(N_f2, N_f1)
46 ...
47 # Print confusion matrix
48 ids = phoneme_id[np.isin(phoneme_id, [1, 2])]
49 X1 = X[ids == 1]
50 X2 = X[ids == 2]
51 plt.scatter(X1[:, 0] - min_f1, X1[:, 1] - min_f2, marker='.',
52             ↪ color='red', label='Phoneme 1')
53 plt.scatter(X2[:, 0] - min_f1, X2[:, 1] - min_f2, marker='.',
54             ↪ color='green', label='Phoneme 2')

```

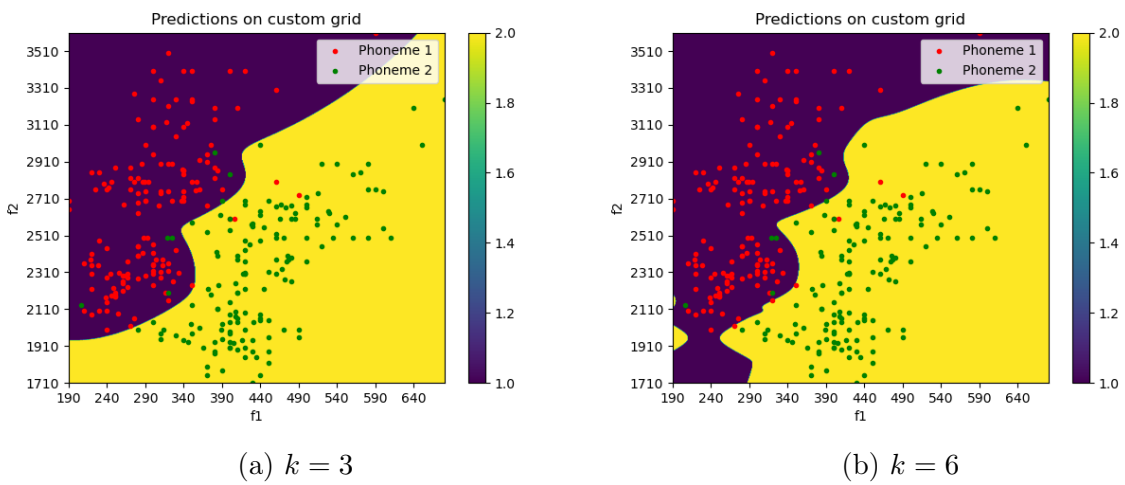


Figure 6: Confusion matrix for k components

The above code is structured as:

1. **Line 2–4:** Copy the relevant frequency values.
2. **Line 6:** Fill `X_phonemes_1_2` with the samples of `X_full` that belong to the chosen phonemes.
3. **Line 9–16:** Create the grid using linearly spaced vectors.
4. **Line 19–45:** Same as explained in [Task 3: Observations](#).
5. **Line 45:** Create `M` array containing 0 in the points that belong to phoneme 1, and 1 in the points that belong to phoneme 2. `M` forms the background for the confusion matrices highlighted in Fig. 6.
6. **Line 51–55:** Print the confusion matrix after sorting the values for the colour assignment.

Task 5

```
1      # Store f1 in the first column, f2 in the second column, and (f1 + f2) in
    ↪   the third column of X_full
2      X_full[:, 0] = f1
3      X_full[:, 1] = f2
4      X_full[:, 2] = f1 + f2
5      ...
6      # Create array containing only samples that belong to phoneme 1
7      X_phoneme_1 = X_full[phoneme_id == p_id, :]
8      ...
9      # Add regularisation to offset singularity
10     s[1, :, :] += 0.001 * np.identity(D)
11
```

Errors

```
Iteration 110/150
Iteration 111/150
Iteration 112/150
C:\Users\user\lib\site-packages\matplotlib\cbook\__init__.py:1333:
  ComplexWarning: Casting complex values to real discards the
    imaginary part
return np.asarray(x, float)
C:\Users\user\lib\site-packages\matplotlib\cbook\__init__.py:1333:
  \linebreak ComplexWarning: Casting complex values to real discards
    the imaginary part
return np.asarray(x, float)
C:\Users\user\lib\site-packages\matplotlib\cbook\__init__.py:1333:
  ComplexWarning: Casting complex values to real discards the
    imaginary part
return np.asarray(x, float)
C:\Users\user\lib\site-packages\matplotlib\cbook\__init__.py:1333:
  ComplexWarning: Casting complex values to real discards the
    imaginary part
return np.asarray(x, float)
Iteration 113/150
Iteration 114/150
```

$$k = 3$$




Figure 7: Full covariance matrices fit to new data

Implemented GMM | Mean values

```
[ 320.95685 2779.7197 3100.6765 ]
[ 290.86166 2585.0154 2875.8772 ]
[ 400.65002 3176.3762 3577.0261 ]
```

Implemented GMM | Covariances

```
[[ 4614.61035882 -445.20477564 4169.40558318]
 [ -445.20477564 6560.32199013 6115.11721449]
 [ 4169.40558318 6115.11721449 10284.52279768]]
[[ 1996.20052277 10892.01222129 12888.21274406]
 [ 10892.01222129 161071.73613414 171963.74835543]
 [ 12888.21274406 171963.74835543 184851.96109951]]
[[ 16314.08843417 37052.91712156 53367.00555573]
 [ 37052.91712156 84205.97468187 121258.89180342]
 [ 53367.00555573 121258.89180342 174625.89735916]]
```

Implemented GMM | Weights

```
[0.22433881 0.74241566 0.03324552]
```

$$k = 6$$




Figure 8: Full covariance matrices fit to new data

Implemented GMM | Mean values

```
[ 303.02448 2678.7605 2981.785 ]
[ 280. 2400. 2680.]
[ 472.704 2742.6934 3215.3975]
[ 268.2653 2047.8208 2316.0862]
[ 251.20546 2205.8057 2457.0112 ]
[ 310.13544 2312.8076 2622.943 ]
```

Implemented GMM | Covariances

```
[[ 3588.25985186 11809.41486645 15397.67471831]
 [ 11809.41486645 137364.33602465 149173.75089109]
 [ 15397.67471831 149173.75089109 164571.42560939]]
[[0.005 0. 0. ]
 [0. 0.005 0. ]
 [0. 0. 0.005]]
[[2228.82332945 -682.77261291 1546.05071654]
 [-682.77261291 750.88929805 68.11668514]
 [1546.05071654 68.11668514 1614.16740169]]
[[ 63.45874286 -110.63331313 -47.17457027]
 [-110.63331313 694.31867819 583.68536506]
 [-47.17457027 583.68536506 536.5107948 ]]
[[ 92.83688005 459.37117444 552.20805449]
 [ 459.37117444 2730.52641071 3189.89758515]
```

```
[ 552.20805449 3189.89758515 3742.10563965]]
[[ 106.54427401 -223.97845157 -117.43417756]
 [-223.97845157 3181.81198929 2957.83353772]
 [-117.43417756 2957.83353772 2840.39936016]]
```

```
Implemented GMM | Weights
```

```
[0.92474453 0.01137863 0.00368719 0.00426987 0.0402883 0.01563148]
```

Observations

The errors displayed in [Errors](#) section display the singularity problem. The issue arises from the covariance matrices that were previously using zero as non-diagonal entries. A singular matrix e.g., A does not have an inverse, where $\det(A) = 0$. The program tries to use the inverse of the covariance matrices that can not exist, and therefore kept timing out.

The problem was fixed by implementing regularisation to the covariance matrices highlighted in Line **10** of the [code](#) sample for this task. A diagonal matrix with the equal dimensions is represented as an identity matrix and multiplied with a value of 0.001 to force the diagonal entries to never reach 0.